# Batch optimization in VW via LBFGS 

Miroslav Dudík

12/16/2011

## Outline

- gradient descent and Newton method
- LBFGS
- LBFGS in VW


## Smooth convex unconstrained optimization

Goal: $\min _{\mathbf{w} \in \mathbb{R}^{d}} f(\mathbf{w})$
where $f$ is strongly convex and twice continuously differentiable

## Smooth convex unconstrained optimization

Goal: $\min _{\mathbf{w} \in \mathbb{R}^{d}} f(\mathbf{w})$
where $f$ is strongly convex and twice continuously differentiable

Our objective:

$$
f(\mathbf{w})=\sum_{i=1}^{n} \operatorname{loss}\left(\mathbf{w} ; x_{i}, y_{i}\right)+\frac{\lambda}{2}\|\mathbf{w}\|^{2}
$$

- possibly weighted loss
- regularization can have coordinate-specific scaling (specified by user)


## Warm-up: Gradient descent

- initialize $\mathrm{w}_{0}$
- for $t=1,2, \ldots$ :
move in the direction of the steepest descent

$$
\mathbf{w}_{t+1}=\mathbf{w}_{t}-\eta \nabla f\left(\mathbf{w}_{t}\right)
$$

## Warm-up: Gradient descent

Gradient descent update:

$$
\mathbf{w}_{t+1}=\mathbf{w}_{t}-\eta \nabla f\left(\mathbf{w}_{t}\right)
$$

## Warm-up: Gradient descent

## gradient

Gradient descent update:

$$
\mathbf{w}_{t+1}=\mathbf{w}_{t}-\eta \nabla f\left(\mathbf{w}_{t}\right)
$$

Equivalently:

- approximate

$$
f(\mathbf{w}) \approx f\left(\mathbf{w}_{t}\right)+\mathbf{g}_{t}^{\top}\left(\mathbf{w}_{t}-\mathbf{w}\right)+\frac{1}{2 \eta}\left\|\mathbf{w}_{t}-\mathbf{w}\right\|^{2}
$$

## Warm-up: Gradient descent

## gradient

Gradient descent update:

$$
\mathbf{w}_{t+1}=\mathbf{w}_{t}-\eta \nabla f\left(\mathbf{w}_{t}\right)
$$

Equivalently:

- approximate

$$
f(\mathbf{w}) \approx f\left(\mathbf{w}_{t}\right)+\mathbf{g}_{t}^{\top}\left(\mathbf{w}_{t}-\mathbf{w}\right)+\frac{1}{2 \eta}\left\|\mathbf{w}_{t}-\mathbf{w}\right\|^{2}
$$

- optimize approximation:

$$
\mathbf{w}_{t+1}=\underset{\mathbf{w}}{\operatorname{argmin}}\left(f\left(\mathbf{w}_{t}\right)+\mathbf{g}_{t}^{\top}\left(\mathbf{w}_{t}-\mathbf{w}\right)+\frac{1}{2 \eta}\left\|\mathbf{w}_{t}-\mathbf{w}\right\|^{2}\right)
$$

## Warm-up: Gradient descent

## gradient

Gradient descent update:

$$
\mathbf{w}_{t+1}=\mathbf{w}_{t}-\eta \nabla f\left(\mathbf{w}_{t}\right)
$$

Equivalently:

- approximate

$$
f(\mathbf{w}) \approx f\left(\mathbf{w}_{t}\right)+\mathbf{g}_{t}^{\top}\left(\mathbf{w}_{t}-\mathbf{w}\right)+\frac{1}{2 \eta}\left\|\mathbf{w}_{t}-\mathbf{w}\right\|^{2}
$$

- optimize approximation:

$$
\mathbf{w}_{t+1}=\underset{\mathbf{w}}{\operatorname{argmin}}\left(f\left(\mathbf{w}_{t}\right)+\mathbf{g}_{t}^{\top}\left(\mathbf{w}_{t}-\mathbf{w}\right)+\frac{1}{2 \eta}\left\|\mathbf{w}_{t}-\mathbf{w}\right\|^{2}\right)
$$

Can we replace quadratic term by a tighter approximation?

## Newton method

$$
\mathbf{H}_{t}=\nabla^{2} f\left(\mathbf{w}_{t}\right)
$$

Better approximation

$$
f(\mathbf{w}) \approx f\left(\mathbf{w}_{t}\right)+\mathbf{g}_{t}^{\top}\left(\mathbf{w}_{t}-\mathbf{w}\right)+\frac{1}{2}\left(\mathbf{w}_{t}-\mathbf{w}\right)^{\top} \mathbf{H}_{t}\left(\mathbf{w}_{t}-\mathbf{w}\right)
$$

Update:

$$
\mathbf{w}_{t+1}=\mathbf{w}_{t}-\mathbf{H}_{t}^{-1} \mathbf{g}_{t}
$$

## Newton method

$$
\mathbf{H}_{t}=\nabla^{2} f\left(\mathbf{w}_{t}\right)
$$

Better approximation

$$
f(\mathbf{w}) \approx f\left(\mathbf{w}_{t}\right)+\mathbf{g}_{t}^{\top}\left(\mathbf{w}_{t}-\mathbf{w}\right)+\frac{1}{2}\left(\mathbf{w}_{t}-\mathbf{w}\right)^{\top} \mathbf{H}_{t}\left(\mathbf{w}_{t}-\mathbf{w}\right)
$$

Update:

$$
\mathbf{w}_{t+1}=\mathbf{w}_{t}-\mathbf{H}_{t}^{-1} \mathbf{g}_{t}
$$

Problem: Hessian can be too big (matrix of size $d x d$ )

## LBFGS = a quasi-Newton method

[Nocedal 1980, Liu-Nocedal 1989]
Instead of the Newton update

$$
\mathbf{w}_{t+1}=\mathbf{w}_{t}-\mathbf{H}_{t}^{-1} \mathbf{g}_{t}
$$

Perform a quasi-Newton update:

$$
\mathbf{w}_{t+1}=\mathbf{w}_{t}-\eta_{t} \mathbf{K}_{t} \mathbf{g}_{t}
$$

where: $\mathbf{K}_{t}$ is a low-rank approximation of $\mathbf{H}_{t}^{-1}$ $\eta_{t}$ is obtained by line search

## LBFGS = a quasi-Newton method

[Nocedal 1980, Liu-Nocedal 1989]
Instead of the Newton update

$$
\mathbf{w}_{t+1}=\mathbf{w}_{t}-\mathbf{H}_{t}^{-1} \mathbf{g}_{t}
$$

Perform a quasi-Newton update:

$$
\mathbf{w}_{t+1}=\mathbf{w}_{t}-\eta_{t} \mathbf{K}_{t} \mathbf{g}_{t}
$$

where: $\mathbf{K}_{t}$ is a low-rank approximation of $\mathbf{H}_{t}^{-1}$ $\eta_{t}$ is obtained by line search

- rank $m$ specified by user (default $m=15$ )
- instead of storage $d^{2}$, only storage $2 d m$ required (update of $\mathrm{K}_{t}$ also has running time $\mathrm{O}(\mathrm{dm})$ per iteration)


## Line search in LBFGS

[Nocedal 1980, Liu-Nocedal 1989]
Update:

$$
\mathbf{w}_{t+1}=\mathbf{w}_{t}-\eta_{t} \mathbf{K}_{t} \mathbf{g}_{t}
$$

- direction determined by $\mathrm{K}_{t} \mathbf{g}_{t}$
- step size $\eta_{t}$ must satisfy Wolfe conditions


## $1^{\text {st }}$ Wolfe condition:



## $1^{\text {st }}$ Wolfe condition:



## $1^{\text {st }}$ Wolfe condition:



## $1^{\text {st }}$ Wolfe condition:



## $1^{\text {st }}$ Wolfe condition:

## $f\left(\mathbf{w}_{t+1}\right) \leq f\left(\mathbf{w}_{t}\right)+\alpha \mathbf{g}_{t}^{\top} \Delta \mathbf{w} \quad$ for some $\alpha$ in $(0,0.5)$


change in w

$$
\Delta \mathbf{w}=\mathbf{w}_{t+1}-\mathbf{w}_{t}
$$

$\mathbf{w}_{t}$ $\mathbf{w}_{t+1}$

## $1^{\text {st }}$ Wolfe condition:

$f\left(\mathbf{w}_{t+1}\right) \leq f\left(\mathbf{w}_{t}\right)+\alpha \mathbf{g}_{t}^{\top} \Delta \mathbf{w} \quad$ for some $\alpha$ in $(0,0.5)$
Rewrite as
$\Delta f \leq \alpha \mathbf{g}_{t}^{\top} \Delta \mathbf{w}$
where $\Delta f=f\left(\mathbf{w}_{t+1}\right)-f\left(\mathbf{w}_{t}\right)$

## $1^{\text {st }}$ Wolfe condition:

$f\left(\mathbf{w}_{t+1}\right) \leq f\left(\mathbf{w}_{t}\right)+\alpha \mathbf{g}_{t}^{\top} \Delta \mathbf{w} \quad$ for some $\alpha$ in $(0,0.5)$
Rewrite as
$\Delta f \leq \alpha \mathbf{g}_{t}^{\top} \Delta \mathbf{w}$
where $\Delta f=f\left(\mathbf{w}_{t+1}\right)-f\left(\mathbf{w}_{t}\right)$

Equivalent to: $\alpha \leq \frac{\Delta f}{\mathbf{g}_{t}^{\top} \Delta \mathbf{w}}$
(because $\mathbf{g}_{t}^{\top} \Delta \mathbf{w}$ is negative)
We use notation wolfel $=\frac{\Delta f}{\mathbf{g}_{t}^{\top} \Delta \mathbf{w}}$ for the ratio on the rhs.

## $2^{\text {nd }}$ Wolfe condition (strengthened):



## $2^{\text {nd }}$ Wolfe condition (strengthened):

 $\left|\mathbf{g}_{t+1}^{\top} \Delta \mathbf{w}\right| \leq \beta \mathbf{g}_{t}^{\top} \Delta \mathbf{w}$
$2^{\text {nd }}$ Wolfe condition (strengthened):
$\left|\mathbf{g}_{t+1}^{\top} \Delta \mathbf{w}\right| \leq \beta \mathbf{g}_{t}^{\top} \Delta \mathbf{w} \quad$ for some $\beta$ in $(\alpha, 1)$
Rewrite as $\left.\beta \geq \frac{\mathbf{g}_{t+1}^{\top} \Delta w}{\mathbf{g}_{t}^{\top} \Delta w} \right\rvert\,$.
We use notation wolfe $2=\frac{\mathbf{g}_{t+1}^{\top} \Delta \mathbf{w}}{\mathbf{g}_{t}^{\top} \Delta \mathbf{w}}$ for the ratio on the rhs.

## Summarizing Wolfe conditions

Let wolfe1 $=\frac{\Delta f}{\mathbf{g}_{t}^{\top} \Delta \mathbf{w}}$ and wolfe2 $=\frac{\mathbf{g}_{t+1}^{\top} \Delta \mathbf{w}}{\mathbf{g}_{t}^{\top} \Delta \mathbf{w}}$.
Let $0<\alpha<0.5, \alpha<\beta<1$.
i) wolfe $1 \geq \alpha$
ii) $\mid$ wolfe $2 \mid \leq \beta$

In VW, the Wolfe conditions are not enforced

- ratios wolfe1 and wolfe2 are logged
- it is always possible to choose $\alpha$ and $\beta$ in the hindsight as long as:
wolfe1>0 and -1<wolfe2<1


## Line search and termination in VW

- in the first iteration:
- evaluate directional $2^{\text {nd }}$ derivative and initialize step size according to the one-dimensional Newton step
- if the loss does not decrease (i.e., wolfe1<0), shrink the step
- in the subsequent iterations:
- set step size to 1.0
- if the loss does not decrease (i.e., wolfe $1<0$ ), shrink the step
- terminate if
either: the specified number of passes over the data is reached
or: $\quad$ the relative decrease in the objective $f(\mathbf{w})$ falls below a threshold


## LBFGS switches

--bfgs
turn on LBFGS optimization
--I2 0.0
L2 regularization coefficient
--mem 15
rank of the inverse Hessian approximation
--termination 0.001
termination threshold for the relative loss decrease

